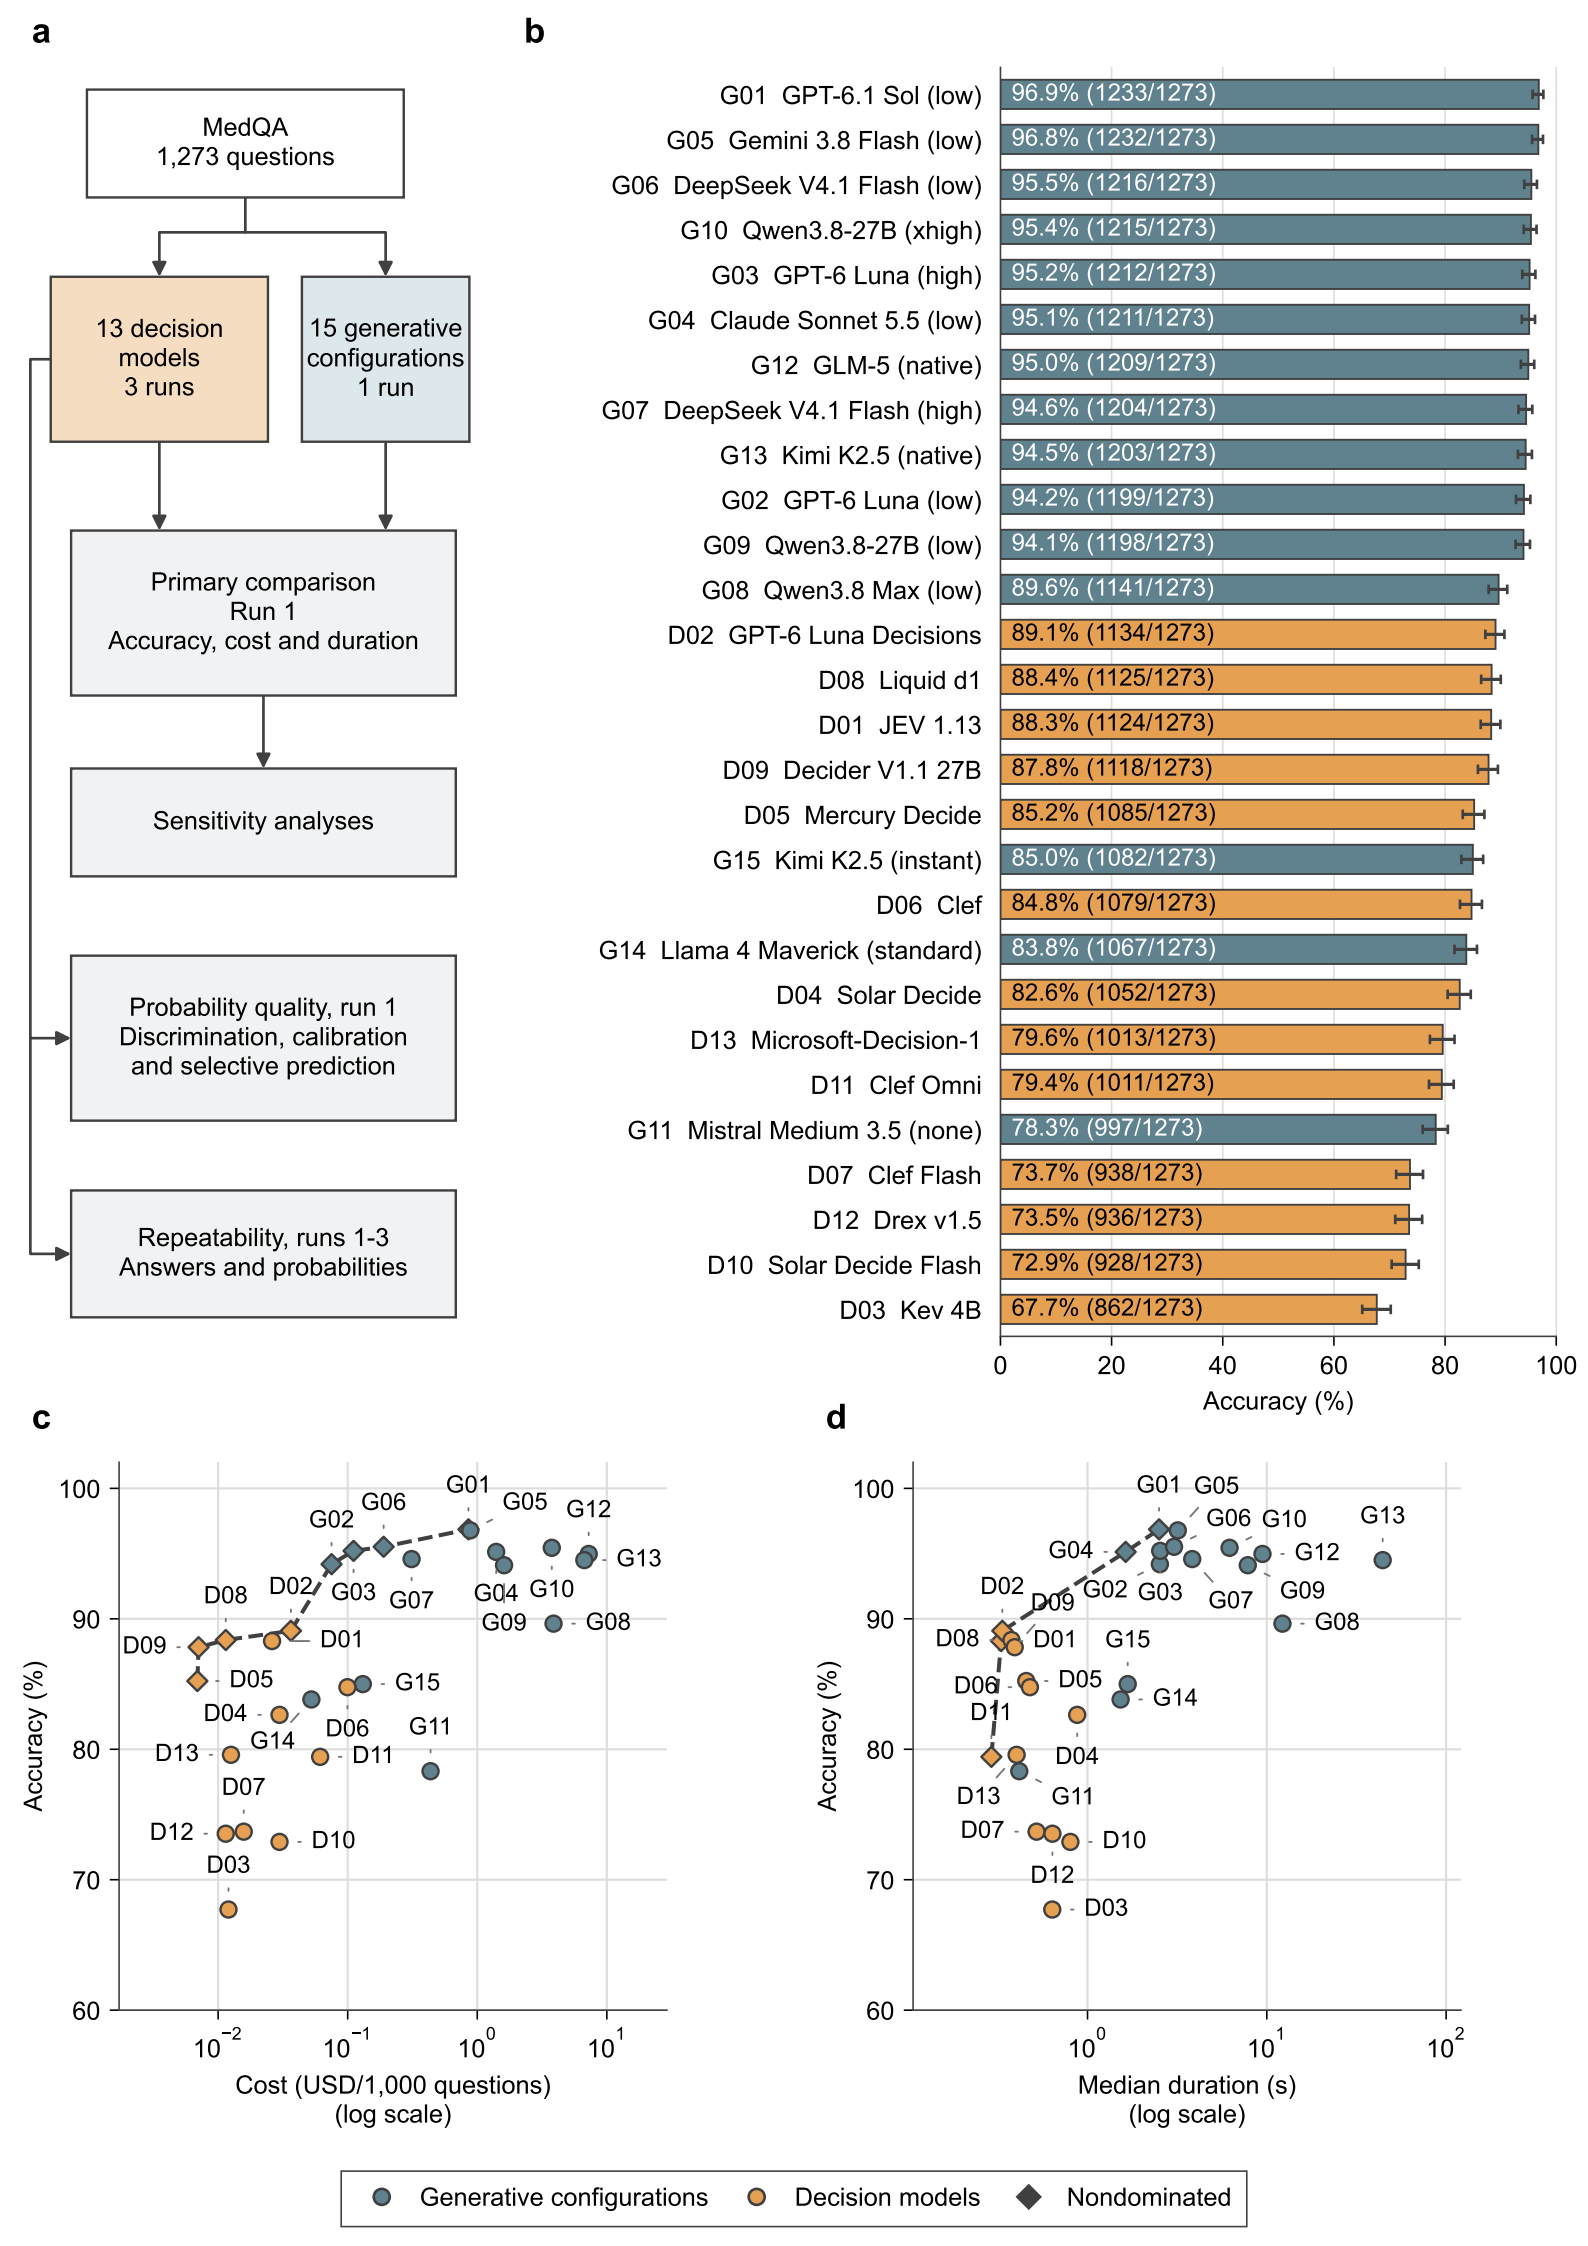
a
MedQA
1,273 questions
13 decision models
3 runs
15 generative configurations
1 run
Primary comparison
Run 1
Accuracy, cost and duration
Sensitivity analyses
Probability quality, run 1
Discrimination, calibration and selective prediction
Repeatability, runs 1-3
Answers and probabilities
b
G01 GPT-6.1 Sol (low) 96.9% (1233/1273)
G05 Gemini 3.8 Flash (low) 96.8% (1232/1273)
G06 DeepSeek V4.1 Flash (low) 95.5% (1216/1273)
G10 Qwen3.8-27B (xhigh) 95.4% (1215/1273)
G03 GPT-6 Luna (high) 95.2% (1212/1273)
G04 Claude Sonnet 5.5 (low) 95.1% (1211/1273)
G12 GLM-5 (native) 95.0% (1209/1273)
G07 DeepSeek V4.1 Flash (high) 94.6% (1204/1273)
G13 Kimi K2.5 (native) 94.5% (1203/1273)
G02 GPT-6 Luna (low) 94.2% (1199/1273)
G09 Qwen3.8-27B (low) 94.1% (1198/1273)
G08 Qwen3.8 Max (low) 89.6% (1141/1273)
D02 GPT-6 Luna Decisions 89.1% (1134/1273)
D08 Liquid d1 88.4% (1125/1273)
D01 JEV 1.13 88.3% (1124/1273)
D09 Decider V1.1 27B 87.8% (1118/1273)
D05 Mercury Decide 85.2% (1085/1273)
G15 Kimi K2.5 (instant) 85.0% (1082/1273)
D06 Clef 84.8% (1079/1273)
G14 Llama 4 Maverick (standard) 83.8% (1067/1273)
D04 Solar Decide 82.6% (1052/1273)
D13 Microsoft-Decision-1 79.6% (1013/1273)
D11 Clef Omni 79.4% (1011/1273)
G11 Mistral Medium 3.5 (none) 78.3% (997/1273)
D07 Clef Flash 73.7% (938/1273)
D12 Drex v1.5 73.5% (936/1273)
D10 Solar Decide Flash 72.9% (928/1273)
D03 Kev 4B 67.7% (862/1273)
0 20 40 60 80 100
Accuracy (%)
c
Accuracy (%)
Cost (USD/1,000 questions) (log scale)
d
Accuracy (%)
Median duration (s) (log scale)
Generative configurations
Decision models
Nondominated